Replay changes sampling and reward scale

Control and experiment

Baseline L2_t1 (d807d91d): temperature1, no replay,128 fresh prompts. BCR16_T1 (73acea80):112 fresh + up to16 replay,400-prompt buffer,worst greedy word score,seed42. Replay regenerates outputs using current policy.

Reward and shared settings

Reward divisor4 →3; same component weights and clip[-2,2]. This scales the same weighted sum by4/3. Both: ReMax,Qwen9B,LoRA320,LR3e-6 constant,8 rollouts,16 GPUs,two epochs; validation temperature0/n1.

Provenance and limits

Baseline03680b38 →target1d72d5a. Scorer,normalizer,actor loss and core advantage code identical; replay integration changes trainer input sampling. Multiple recoveries reset replay; equal steps have different fresh-data exposure.

Training stopped

User stop confirmed by Ray: recovery750/driver10000000 STOPPED,0 allocated GPUs. Monitor deleted. Historical last confirmed optimizer755; recovery750 stopped during initial validation with no new optimizer update.

Pinned source audit and collected W&B configs; live stop verified by exact Ray submission.