Check scoring-sensitive changes before tuning rewards

Evidence

At step600: LS clean−24 errors,LS other−42,Earnings+3,OpenML−2 net errors. OpenML has2 samples with changed n_ref,so its corpus ratio needs both denominators. Table snapshots match20/20 examples per run; the full dumps cover datasets absent from selected tables.

Bounded next check

Pinned scorer replay reproduces all12 selected cases. In OpenML ID1174608741297464509 2,parenthesized(2006) disappears from the normalized hypothesis: errors2 →6. Audit bracket/number normalization first; retain official scores and test a symmetric diagnostic convention on both runs.

Success criteria

Keep official per-dataset p_err unchanged in the report. Any later one-factor trial should fix divisor/exposure confounds,reduce verified scoring-pattern errors at matched budget,and guard LS clean,LS other,OpenML and Earnings separately. No new training was launched.

Pinned source audit and collected W&B configs; live stop verified by exact Ray submission.